


Motivation

Long Running Shift to Concurrency


*For 30 years computer performance driven by **Moore's Law**;
from now on, it will be driven by **Amdahl's Law**.*



*Moore's Law is delivering
more cores but not faster cores.*

March 2006

Amdahl's Law



Theoretical max speedup using **multiple processors?**

Limited by the time needed for **sequential processing.**

Scale Up

- Remember SOA ? EJB 1 ?
 - distributed components, then services
 - **functional boundaries** between system components
 - location transparent resources **pooled**
- ... but it didn't do well
 - complex / expensive solutions
 - full-stack approach
 - overhead not factored in: **latency failure tolerance**
- Ended up in hype failure mockeries and hacker news rants

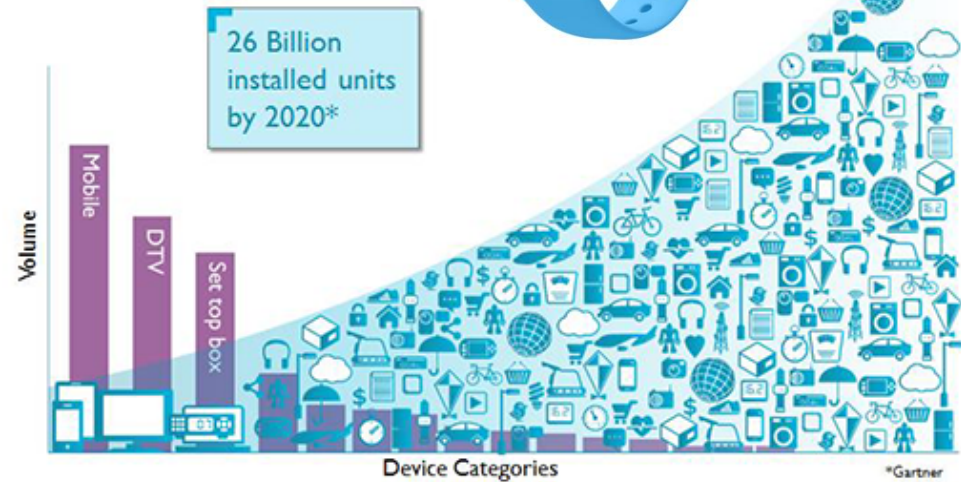
Internet Scale

- Social media became mainstream
 - massive scale
- *Background Processing* and pooled resources regain interest
 - everything “**write-heavy**” -- **candidate for async handling**
 - e.g. a tweet
- Push-based servers start to proliferate
 - async response handling (e.g. facebook messages)

Cloud Scale

- “Cloud-ready” systems began to spread around 2012
 - AWS, open-source PaaS democratization
- Tech giants expose internal distributed stack
 - Twitter, Google, Netflix...
- Tools to address core concerns of distributed systems
 - **concurrency, service discovery, monitoring**

Wait! Another traffic spike: IoT



Scale Out

- There are **limits to scale** the monolith way
 - cost efficiency aside
- **Memory overuse**
 - large thread pools, application footprint
- **CPU underuse**
 - blocking I/O (DB, remote service)

Tales of Time and Space

- Massive scale requires higher efficiency
 - respect physical laws
 - **speed of light dampens inter-service** communication
- Pooling resources simply not enough
 - **need to coordinate, compose, orchestrate data flows**
- **Non-blocking must be embraced** fundamentally

Non-Blocking Architecture

- Non-blocking requires profound change
 - **can't write imperative code**
 - **can't assume single thread** of control (e.g. exception handling)
- Must deal with async results
 - **listener/callbacks become unwieldy** with nested composition
- Everything becomes an event stream
 - e.g. from `InputStream` to **stream of events**

Reactive Programming

Generally speaking...

- **“Reactive”** is used broadly to define event-driven systems
 - UI events, network protocols
- Now also entering the domain of application/business logic
- **Non-blocking** event-driven architecture

Reactive Manifesto

- Reflects the emerging field of scalable, non-blocking applications
- A hark back to other successful manifestos
- Well articulated but very broad strokes
- Defines qualities of **reactive systems**
- For **reactive programming** we need concrete tools

Reactive Programming in Context

- To build reactive apps we need two essential building blocks
- **Contract for interop between non-blocking components**
 - Reactive Streams spec
- **API for composing asynchronous programs**
 - e.g. Reactive Extensions

Myths about “Reactive”

- Async / concurrent == reactive ?
 - easy to end up with too many threads
 - fill up hand-off queue
- Must be async to be reactive ?
 - nope but must be agnostic to source of concurrency
- Reactive is the domain of specific programming languages ?
 - language features can help
 - e.g. JDK 8 lambda

Reactive Streams Specification for the JVM

Specification Goals

- Govern the exchange of data **across async boundaries**
- **Use back-pressure** for flow control
- Fast sources should not overwhelm slower consumers
- Provide **API** types, **TCK**

Back-Pressure

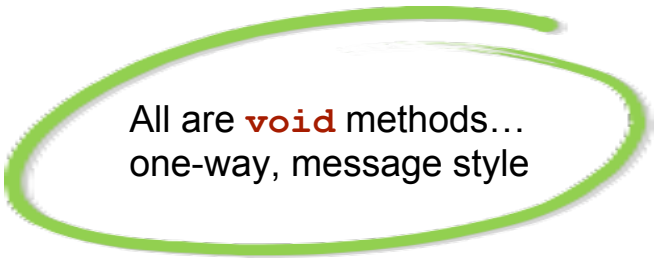
- When publisher maintains **higher rate for extended time**
 - **queues grow** without bounds
- Back-pressure allows a subscriber to **control queue bounds**
 - subscriber **requests #** of elements
 - publisher produces **up to #** of requested
- If source can't be controlled (e.g. mouse movement clock tick)
 - publisher may buffer or drop
 - must obey **#** of requested elements

API Types

```
public interface Publisher<T> {  
    void subscribe(Subscriber<? super T> s);  
}
```

```
public interface Subscriber<T> {  
    void onSubscribe(Subscription s);  
    void onNext(T t);  
    void onError(Throwable t);  
    void onComplete();  
}
```

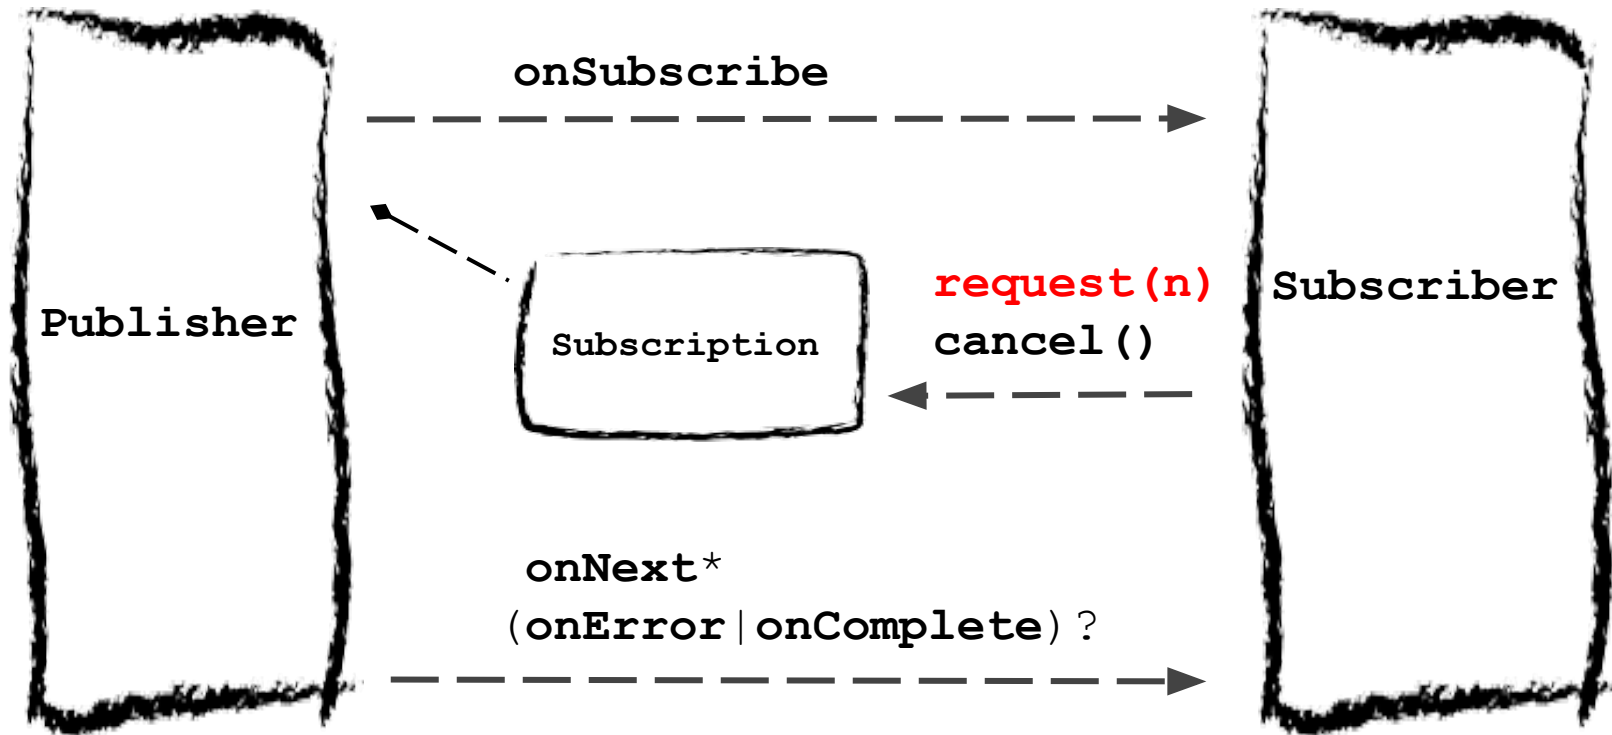
```
public interface Subscription {  
    void request(long n);  
    void cancel();  
}
```




All are **void** methods...
one-way, message style

```
public interface Processor<T R> extends Subscriber<T> Publisher<R> {  
}
```

Reactive Streams API




Reactive Streams: Message Passing



onSubscribe(Subscription)




subscribe(Subscriber)




onSubscribe(Subscription)







onSubscribe(Subscription)








onSubscribe(Subscription)




















Spec Rules

- **No concurrent** calls from Publisher to Subscriber (1.3) or vice versa (2.7)
- A Subscriber may perform work **synchronously or asynchronously**
 - either way **must be non-blocking** (2.2)
- Upper bound on open recursion required (3.3)
 - `onNext` → `request` → `onNext` → ... → `StackOverflow`
- No exceptions can be raised except NPE for `null` input (1.9 2.13)
 - Publisher calls `onError(Throwable)`
 - Subscriber cancels subscription

What does request (n) mean?

- N represents an element count
 - relative (byte chunks) more so than absolute (# of bytes)
- Boundaries
 - Must be > 0
 - `Long.MAX_VALUE` means unbounded read
 - So does overflow of N

Request Strategies

- Request some receive all request more ... (“*stop-and-wait*”)
 - request signal **equivalent to an ack**
- Pre-fill buffer of certain capacity (“*pre-fetch*”)
 - request again **when buffer falls** below some threshold
 - more in-flight data → more risk
- Request using timer/**scheduler** facility (“*poll*”)
- Adapt to network or processing latency (“*dynamic*”)
 - find and pick the **best request size**

Demo:

Publisher & Subscriber, Back-pressure, TCK verification
[Demo source](#), [TCK test](#)

Async vs Sync

- Sync is explicitly allowed -- 2.2 3.2 3.10 3.11
- If `Publisher/Subscriber` are both sync → open recursion
 - not always an issue e.g. single item `Publisher`
- The picture is different in a processing chain (vs. single Publisher-Subscriber)
 - async hand-off is likely needed at some stage
 - but not at every stage
- Async is not the goal non-blocking is the overall objective

Async vs Sync (continued)

- Implementations are free to choose how to do this
- As long as they comply with all the rules
- Generally speaking the burden is on the `Publisher`
 - make async or sync calls as necessary
 - prevent open recursion
 - sequential signals
- *Subscribers can largely ignore such concerns*

Reactive Streams → JDK 9 Flow.java

- No single best fluent async/parallel API in Java ^[1]
- `CompletableFuture/CompletionStage`
 - continuation-style programming on futures
- `java.util.stream`
 - multi-stage/parallel, “pull”-style operations on collections
- No “push”-style API for items as they become available from active source

Java “concurrency-interest” list:

[1] [initial announcement](#) [2] [why in the JDK?](#) [3] [recent update](#)

JDK 9 `java.util.concurrent`

- [`java.util.concurrent.Flow`](#)
 - `Flow.Publisher` `Flow.Subscriber` `Flow.Subscription` ...
- [`Flow.consume\(...\)`](#) and [`Flow.stream\(...\)`](#)
 - Tie-ins to `CompletableFuture` and `java.util.Stream`
- [`SubmissionPublisher`](#)
 - turn any kind of source to `Publisher`
 - `Executor`-based delivery to subscribers w/ bounded (ring) buffer
 - `submit` and `offer` strategies (block or drop) to publish

API For Async Composition

Application


Application

Reactive
Streams

HTTP

Data

Broker



API to **compose**
asynchronous
programs?

Working with Streams

- Non-blocking services return `Publisher<T>` instead of `<T>`
- How do you attach further processing?
- Reactive Streams is a callback-based API
 - becomes very nested quickly
- Need something more declarative
 - it's beyond the scope of Reactive Streams

Stream Operations

- `Publisher` represents a stream of data
- It's natural to apply operations functional-style
 - like the Java 8 Stream
- Need API for composing async logic
 - rise above callbacks

Reactor Stream

- [Project Reactor](#) provides a Stream API
- Reactive Streams `Publisher` + composable operations

```
Streams.just('a' 'b' 'c')  
    .take(2)  
    .map(Character::toUpperCase)  
    .consume(System.out::println);
```

Demo:

Reactor Stream, Back-Pressure

[Demo source](#)

Reactive Extensions (Rx)

- Stream composition API based on Observer pattern
- Originated at Microsoft, work by Erik Meijer
- Implemented for different languages -- RxJava, RxJS, Rx.NET, ...

```
Observable.just('a', 'b', 'c')  
    .take(2)  
    .map(Character::toUpperCase)  
    .subscribe(System.out::println);
```

RxJava and Reactive Streams

- RxJava 1.x predates Reactive Streams and doesn't implement it directly
- Very similar concepts, different names
 - `Observable-Observer` vs `Publisher-Subscriber`
- RxJava supports “reactive pull” back-pressure
- [RxJava 1.x - Reactive Streams](#) bridge

Rx vs Java 8 Streams

- Java 8 Streams best suited for collections
- `Observer` and `Iterator` actually very closely related
 - same except push vs pull
- `Observable` however can represent any source including collections
- RxJava 2 is planned to support Java 8 Stream
 - `Observable.from(Iterable)` + implement `java.util.Stream`

Demo:

RxJava, Back-Pressure, Reactive Streams Bridge

[Demo class](#), [TCK test](#)

RxJava vs Reactor

- RxJava is a great choice for composition in applications
 - many operators, polyglot (client/server), well-documented
- Reactive Streams is ideal for use in library APIs
 - remain agnostic to composition
- Reactor is positioned as foundation for libraries
 - essential Reactive Streams infrastructure + core operators
 - also used in high volume applications

Head Start with Stream Composition

Operators: reactivex.io Alphabetical List

Aggregate, All, Amb, and_, And, Any, apply, as_blocking, AsObservable, AssertEqual, asyncAction, asyncFunc, Average, averageDouble, averageFloat, averageInteger, averageLong, blocking, Buffer, bufferWithCount, bufferWithTime, bufferWithTimeOrCount, byLine, cache, case, Cast, Catch, catchException, collect, collect (RxScala version of Filter), CombineLatest, combineLatestWith, Concat, concat_all, concatMap, concatMapObserver, concatAll, concatWith, Connect, connect_forever, cons, Contains, controlled, Count, countLong, Create, cycle, Debounce, decode, DefaultIfEmpty, Defer, deferFuture, Delay, delaySubscription, delayWithSelector, Dematerialize, Distinct, DistinctUntilChanged, Do, doAction, doOnCompleted, doOnError, doOnRequest, doOnSubscribe, doOnTerminate, doOnUnsubscribe, doseq, doWhile, drop, dropRight, dropUntil, dropWhile, ElementAt, ElementAtOrDefault, Empty, empty?, encode, ensures, error, every, exclusive, exists, expand, failWith, Filter, filterNot, Finally, finallyAction, finallyDo, find, findIndex, First, FirstOrDefault, firstOrElse, FlatMap, flatMapFirst, flatMapIterable, flatMapIterableWith, flatMapLatest, flatMapObserver, flatMapWith, flatMapWithMaxConcurrent, flat_map_with_index, flatten, flattenDelayError, foldl, foldLeft, for, forall, ForEach, forEachFuture, forIn, forkJoin, From, fromAction, fromArray, FromAsyncPattern, fromCallable, fromCallback, FromEvent, FromEventPattern, fromFunc0, from_future, from_iterable, from_list, fromNodeCallback, fromPromise, fromRunnable, Generate, generateWithAbsoluteTime, generateWithRelativeTime, generator, GetEnumerator, getIterator, GroupBy, GroupByUntil, GroupJoin, head, headOption, headOrElse, if, ifThen, IgnoreElements, indexOf, interleave, interpose, Interval, into, isEmpty, items, Join, join (string), jortSort, jortSortUntil, Just, keep, keep-indexed, Last, lastOption, LastOrDefault, lastOrElse, Latest, latest (Rx.rb version of Switch), length, let, letBind, limit, LongCount, ManySelect, Map, map (RxClojure version of Zip), MapCat, mapCat (RxClojure version of Zip), map-indexed, map_with_index, Materialize, Max, MaxBy, Merge, mergeAll, merge_concurrent, mergeDelayError, mergeObservable, mergeWith, Min, MinBy, MostRecent, Multicast, nest, Never, Next, Next (BlockingObservable version), none, nonEmpty, nth, ObserveOn, ObserveOnDispatcher, observeSingleOn, of, of_array, ofArrayChanges, of_enumerable, of_enumerator, ofObjectChanges, OfType, ofWithScheduler, onBackpressureBlock, onBackpressureBuffer, onBackpressureDrop, OnErrorResumeNext, onErrorReturn, onExceptionResumeNext, orElse, pairs, pairwise, partition, partition-all, pausable, pausableBuffered, pluck, product, Publish, PublishLast, publish_synchronized, publishValue, raise_error, Range, Reduce, reductions, RefCount, Repeat, repeat_indefinitely, repeatWhen, Replay, rescue_error, rest, Retry, retry_indefinitely, retryWhen, Return, returnElement, returnValue, runAsync, Sample, Scan, scope, Select (alternate name of Map), select (alternate name of Filter), selectConcat, selectConcatObserver, SelectMany, selectManyObserver, select switch, selectSwitch, selectSwitchFirst, selectWithMaxConcurrent, select_with_index, seq, SequenceEqual, sequence_eql?, SequenceEqualWith, Serialize, share, shareReplay, shareValue, Single, SingleOrDefault, singleOption, singleOrElse, size, Skip, SkipLast, skipLastWithTime, SkipUntil, skipUntilWithTime, SkipWhile, skip_while_with_index, skip_with_time, slice, sliding, slidingBuffer, some, sort, sort-by, sorted-list-by, split, split-with, Start, startAsync, startFuture, StartWith, stringConcat, stopAndWait, subscribe, SubscribeOn, SubscribeOnDispatcher, subscribeOnCompleted, subscribeOnError, subscribeOnNext, Sum, sumDouble, sumFloat, sumInteger, sumLong, Switch, switchCase, switchIfEmpty, switchLatest, switchMap, switchOnNext, Synchronize, Take, take_with_time, takeFirst, TakeLast, takeLastBuffer, takeLastBufferWithTime, takeLastWithTime, takeRight (see also: TakeLast), TakeUntil, takeUntilWithTime, TakeWhile, take_while_with_index, tail, tap, tapOnCompleted, tapOnError, tapOnNext, Then, thenDo, Throttle, throttleFirst, throttleLast, throttleWithSelector, throttleWithTimeout, Throw, throwError, throwException, TimeInterval, Timeout, timeoutWithSelector, Timer, Timestamp, To, to_a, ToArray, ToAsync, toBlocking, toBuffer, to_dict, ToDictionary, ToEnumerable, ToEvent, ToEventPattern, ToFuture, to_h, toIndexedSeq, toIterable, toIterator, ToList, ToLookup, toMap, toMultiMap, ToObservable, toSet, toSortedList, toStream, ToTask, toTraversable, toVector, tumbling, tumblingBuffer, unsubscribeOn, Using, When, Where, while, whileDo, Window, windowWithCount, windowWithTime, windowWithTimeOrCount, windowed, withFilter, withLatestFrom, Zip, zipArray, zipWith, zipWithIndex

Operators: rx.Observable

all, amb, ambWith, asObservable, buffer, cache, cast, collect, combineLatest, compose, concat, concatMap, concatWith, contains, count, countLong, create, debounce, defaultIfEmpty, defer, delay, delaySubscription, dematerialize, distinct, distinctUntilChanged, doOnCompleted, doOnEach, doOnError, doOnNext, doOnRequest, doOnSubscribe, doOnTerminate, doOnUnsubscribe, elementAt, elementAtOrDefault, empty, error, exists, filter, finallyDo, first, firstOrDefault, flatMap, flatMapIterable, from, groupBy, groupJoin, ignoreElements, interval, isEmpty, join, just, last, lastOrDefault, lift, limit, map, materialize, merge, mergeDelayError, mergeWith, nest, never, observeOn, ofType, onBackpressureBlock, onBackpressureBuffer, onBackpressureDrop, onBackpressureLatest, onErrorResumeNext, onErrorReturn, onExceptionResumeNext, publish, range, reduce, repeat, repeatWhen, replay, retry, retryWhen, sample, scan, sequenceEqual, serialize, share, single, singleOrDefault, skip, skipLast, skipUntil, skipWhile, startWith, subscribeOn, switchIfEmpty, switchMap, switchOnNext, take, takeFirst, takeLast, takeLastBuffer, takeUntil, takeWhile, throttleFirst, throttleLast, throttleWithTimeout, timeInterval, timeout, timer, timestamp, toList, toMap, toMultimap, toSortedList, unsubscribeOn, using, window, withLatestFrom, zip, zipWith

Operators: what can you do with sequences?

- Originate -- fixed values, arrays, ranges, from scratch
- Reduce -- filter, accumulate, aggregate, partition
- Transform -- create new type of sequence
- Combine -- concatenate, merge, pair
- Control -- overflow, errors

Compose Non-blocking Service Layer

Blocking:

```
List<String> findUsers(String skill);  
LinkedInProfile getConnections(String id);  
TwitterProfile getTwitterProfile(String id);  
FacebookProfile getFacebookProfile(String id);
```

Non-Blocking:

```
Observable<String> findUsers(String skill);  
Observable<LinkedInProfile> getConnections(String id);  
Observable<TwitterProfile> getTwitterProfile(String id);  
Observable<FacebookProfile> getFacebookProfile(String id);
```

Demo:

Get List of RxJava Operators

[Demo source](#)

Compose Non-Blocking Service Layer

[Demo source](#)

flatMap?!



From map to flatMap

- map is used to apply a function to each element
 - `Function<T, R>`
- The function could return a new sequence
 - `Function<T, Observable<R>>`
- In which case we go from **one** to **many** sequences
- flatMap merges (flattens) those back into one
 - i.e. map + merge

Learn Your Operators

- `flatMap` is essential composing non-blocking services
 - scatter-gather, microservices
- Along with `concatMap`, `merge`, `zip`, and others
- Learn and experiment with operators
- Use [decision tree](#) for choosing operators on reactivex.io to learn

Stream Types

- **“Cold”** source -- **passive**, subscriber can dictate rate
 - e.g. file, database cursor
 - usually “replayed from the top” per subscriber
- **“Hot”** source -- **active**, produce irrespective of subscribers
 - e.g. mouse events, timer events
 - cannot repeat
- Not always a clear distinction
 - operators further change stream behavior
 - e.g. merge hot + cold

Stream Types & Back-Pressure

- Cold sources are well suited for reactive back-pressure
 - i.e. support `request(n)`
- Hot sources cannot pause, need alternative flow control
 - buffer, sample or drop
- Several operators exist
 - `onBackPressureBuffer`, `onBackPressureDrop`
 - `buffer`, `window`
 - etc.

More on Stream Composition

How Operators Work

- Each operator is a **deferred** declaration of work to be done

```
observable.map(...).filter(...).take(5) ...
```

- No work is done until there is a subscriber

```
observable.map(...).filter(...).take(5) subscribe(...)
```

- Underlying this is the `lift` operator

How the `lift` Operator Works

- Accepts function for decorating `Subscriber`
 - `lift(Function<Subscriber<DOWN>, Subscriber<UP>>)`
- Returns new `Observable` that when subscribed to will use the function to decorate the `Subscriber`
- Effectively each operator decorates the target `Subscriber`
 - at the time of subscribing

The Subscribe

- For every subscriber the `lift` chain is used (subscriber decorated)
- Data starts flowing to the subscriber
- Every subscriber is decorated independently
- Results in separate data pipelines (N subscribers → N pipelines)

Reasons For Multiple Subscribers

- Different rates of consumption, slow/fast consumers
- Different resource needs, i.e. IO vs CPU bound
- Different tolerance to errors
- Different views on dealing with overflow
- Different code bases

Shared Data Pipeline

- It's possible to have shared pipeline across subscribers
- Insert shared publisher in the chain
 - vs `lift` which decorates every subscriber
- Both RxJava and Reactor support this
 - `observable.share()` and `stream.process(...)` respectively
- Reactor has processors for fan-out or point-to-point
 - pass all elements to all subscribers
 - distribute elements (“worker queue” pattern)

Demo:

*Multiple Streams,
Shared Streams (fan-out, point-to-point)*

[Demo Source](#)

Concurrency in Rx

- By default all work performed on subscriber's thread
- Operators generally not concurrent
 - some do accept `rx.Scheduler` (timer-related)
- Two operators for explicit concurrency control
 - `subscribeOn(Scheduler)`, `observeOn(Scheduler)`
- `Schedulers` class with factory methods for `Scheduler` instances
 - `io()`, `computation()`, `trampoline()`, ...

observeOn Operator

- Invoke rest of the chain of operators on specified Scheduler
- Inserts async boundary for downstream processing
- Uses a queue to buffer elements
- Temporarily mitigate different upstream/downstream rates

subscribeOn operator

- Put source Observable on specified Scheduler
 - both initial subscribe + subsequent requests
- Order in chain of operators not significant
- Use for Observable that performs blocking I/O
 - e.g. `observable.subscribeOn(Schedulers.io())`

Concurrency in Reactor

- Reactive Streams `Processor` for explicit concurrency control
 - comparable to Rx `observeOn`

```
Streams.period(1)
    .process(RingBufferProcessor.create())
    .consume(...);
```

- `Processor` can also insert `async` boundary for upstream signals
 - comparable to Rx `subscribeOn`
- Internally some operators may use `Timer` thread
 - the main reason for Rx operators accepting `Scheduler`

Error Handling

- `try-catch-finally` is of limited value in non-blocking app
- Exceptions may occur on different thread(s)
- Notifications are the only path to consistent error handling
- Reactive Streams forbids propagating exceptions through the call-stack
- Subscriber receives `onError`

Try but won't catch

```
try {  
    Streams.just(1, 2, 3, 4, 5)  
        .map(i -> {  
            throw new NullPointerException();  
        })  
        .consume(  
            element -> { ... },  
            error -> logger.error("Oooh error!", error)  
        );  
}  
catch (Throwable ex) {  
    logger.error("Crickets...");  
}
```

Recovering from Errors

- There are operators to handle error notifications
- Swallow error + emit backup sequence (or single value)
 - `onErrorResumeNext`, `onErrorReturn`
- Re-subscribe, it might work next time
 - `retry`, `retryWhen`
- Lifecycle related
 - `doOnTerminate/finallyDo...` before/after error-complete



SPRINGONE2GX

WASHINGTON, DC

Learn More. Stay Connected.



@springcentral



Spring.io/video

Follow-up talk on Thursday (10:30am)

“Reactive Web Applications”